AI Passport

Useful AI tasks with selected memory and connected evidence.

Live: passport.ajnasnb.com
Public source: github.com/AjnasNB/app-passport
Evaluation date: 2026-10-13

Evidence	Measured result
Real hosted AI and gate cases	8 / 8
Pinned real tool artifacts / controlled MCP calls	10 / 10
Local SDK, graph, API and orchestration checks	71 / 71
Independent graph MCP path / truncation checks	4 / 4
Graph structure	27 nodes, 43 edges
Current signed graph entities	71

The coding handoff generates a bounded release guard and passes six deterministic boundary cases. India number tutor calculates exact paise and lakh/crore values. UPI Safety Quest is a playable three-round learning game. The AI explanations use real Cloudflare-hosted Llama 3.2 3B inference. Hosted model weights are not attested.

These are project-run demonstrations on public or synthetic inputs. They do not establish customer adoption, an independent security audit, universal model correctness, or proof of inference computation. Both demonstration wallets are project controlled.

The working flow

Select a task

Choose coding, number learning or the safety game. No wallet or account is required.

Read current evidence

Read the signed graph from Tiramisu, verify complete manifest membership, accepted authors and entity-index bindings.

Bind the original sources

Read actual public source bytes, check their digests, validate the Qarinah packet and recompute its selected repository snapshot.

Check current admission

Require current accepted context and runtime reviews. Withdrawn, expired, modified or stale evidence blocks the model call.

Run bounded tools and AI

Run deterministic arithmetic or scoring, select compact cited memory, reserve a free demo allowance and call the hosted model.

Validate and trace the output

Generated guards are interpreted and tested, never executed as arbitrary code. Numeric statements come only from the calculator. Inspect source links and download a receipt.

A Hindi model test produced a wrong conversion during development. The corrected flow renders deterministic numeric results separately and withholds model commentary containing unchecked numerical claims. Provenance alone does not guarantee factual accuracy.

Live AI evaluations

Task / condition	HTTP	Inference	Result
coding / current	200	Yes	Passed
numbers / en	200	Yes	Passed
safety / en	200	Yes	Passed
coding / withdrawn	409	No	Passed
coding / expired	409	No	Passed
coding / tampered	409	No	Passed
coding / stale	409	No	Passed
numbers / hi	200	Yes	Passed

Coding memory selection

Selected memory: 1,371 characters. Full Qarinah packet: 10,096 characters. Selected decision events: 2. This compares payload sizes; token estimates use chars/4 and are not an exact tokenizer benchmark.

Tested public tool artifacts

Artifact	Pinned version	Protocol call
Qarinah	0.6.0+1aa7022	Passed
MCP Filesystem	1.0.0	Passed
MCP Memory	1.0.0	Passed
MCP Sequential Thinking	1.0.0	Passed
MCP Everything	1.0.0	Passed
Playwright MCP	0.0.83	Passed
Chrome DevTools MCP	1.10.1	Passed
MCP Git	2026.10.10	Passed
MCP Fetch	2026.10.10	Passed
MCP Time	2026.10.10	Passed

Why Arkiv and how to reproduce

Arkiv stores shared, attributable, queryable evidence with creator and expiry metadata. Another reader can retrieve current review and withdrawal state through its own provider, without this website backend. Consumers choose their accepted issuers and reviewers. A conventional database remains sufficient for a single operator’s calculator or chat; Arkiv is useful here for interoperable evidence, not for doing arithmetic or inference.

The graph is an application layer over Arkiv entities. Nodes, edges and a complete root manifest are separately signed. Traversal runs in the SDK, not native Cypher. Storage is 30 days and signed validity is 14 days, so the operator must renew evidence. RPC chain metadata and query completeness remain provider dependencies.

Independent readers

Install /downloads/app-passport-sdk-0.2.1.tgz. Read /data/policy.json and /data/graph-config.json, then use readGraph(connection, graphDigest) and traverseGraph(graph, "app:coding", {depth:4}). The included MCP bridge exposes trace_evidence through stdio. Its tested consumer uses the raw Tiramisu RPC, without the website or operator quota secret.

Project reproduction

npm ci --ignore-scripts
npm run build
npm run typecheck
npm test
npx tsx scripts/evaluate-ai.ts
npx tsx scripts/test-graph-mcp.ts

Free demo boundaries

At most forty model attempts per UTC day, with three run requests per minute per IP. Source memory is limited to 10,000 characters and output to 420 model tokens. The calculator and game work without inference. Readers require no Web3 tokens; writers use free Tiramisu test tokens. No mainnet transaction or paid plan upgrade is required by this demo.

Primary references

Arkiv queries · Actual model · NPCI guidance
AI receipts · Signed chain evidence · Complete architecture